

Mock Paper 2B — Algorithms & Programming (H446/02)

Time allowed: 2 hours 30 minutes · Total: 140 marks · Answer all questions.

Instructions

- Use **OCR Exam Reference Language** or a **high-level language** for all code.
 - **Show all traces** in full — give the complete variable / stack / table state at every step, not just the final answer.
 - Quality of written communication is assessed in the extended-response question.
 - Mark your work against the **mark scheme** at the end of this paper.
 - Each part is tagged with its mark total [n] and the dominant assessment objective (A01 / A02 / A03) .
-

Questions

Question 1 — Computational thinking applied to a scenario [11]

A logistics company runs **GridStore**, an automated warehouse. Hundreds of robots collect items from shelves and bring them to packing stations. A central controller plans each robot's route, avoids collisions, and reorders jobs when an urgent order arrives.

- (a) The design team models the warehouse floor as a grid of square cells, each marked as *free* or *blocked*, and ignores the physical height of shelving and the colour of the robots. Name the computational thinking technique being used and explain, with reference to GridStore, **why** it is helpful here. [3] (AO2)
- (b) The job of "fulfil one customer order" is broken down into smaller tasks. Identify **two** sensible sub-problems of fulfilling an order and, for each, state one input and one output. [4] (AO2)
- (c) Two robots simultaneously plan routes that pass through the same cell at the same time-step. Explain why this is an example of a problem requiring **concurrency** to be managed, and describe **one** consequence if it is not managed. [2] (AO2)
- (d) The controller must guarantee a fair, predictable response time when reordering jobs. State **one** reason why **abstraction by generalisation (categorisation)** of jobs into "standard" and "urgent" types helps the controller meet this requirement. [2] (AO2)

Running total: 11

Question 2 — Programming techniques [12]

A developer is writing utility subroutines for GridStore.

- (a) State **one** difference between an **iterative** solution and a **recursive** solution to the same problem. [1] (AO1)
- (b) The subroutine below is written in OCR Exam Reference Language. Global variable `count` is declared outside the subroutine.

```

count = 0

procedure adjust(byRef arr, byVal step)
    for i = 0 to arr.length - 1
        arr[i] = arr[i] + step
        count = count + 1
    next i
endprocedure

nums = [10, 20, 30]
adjust(nums, 5)
print(nums)
print(count)

```

State exactly what is printed, in order, and explain your answer with reference to **pass by reference**, **pass by value** and **variable scope**. [5] (AO2)

(c) Explain what is meant by the **scope** of a variable and give **one** disadvantage of relying on a **global** variable such as `count` above. [3] (AO2)

(d) Rewrite the loop below so that it has identical behaviour but uses a **condition-controlled (while)** loop instead of a count-controlled loop: [3] (AO2)

```

total = 0
for k = 1 to 5
    total = total + k
next k

```

Running total: 23

Question 3 — Recursion and the call stack [10]

The recursive function `power` is defined below.

```

function power(base, exp)
    if exp == 0 then
        return 1
    else
        return base * power(base, exp - 1)
    endif
endfunction

```

(a) Identify the **base case** and the **recursive (general) case** of this function. [2] (AO1)

(b) The call `power(3, 4)` is made. Using a **call-stack** trace, show each frame pushed onto the stack and the value returned as each frame is popped (unwinds). [5] (AO2)

(c) State the **final value** returned by `power(3, 4)`. [1] (AO2)

(d) A colleague suggests removing the `if exp == 0` base case so that the function "just keeps multiplying". Explain what happens at run-time if the base case is removed, and name the resulting error. [2] (AO2)

Running total: 33

Question 4 — Object-oriented programming [14]

GridStore models its robots with classes. Write/complete the definitions below (OCR Exam Reference Language or a high-level language).

(a) Define a class `Robot` with **private** attributes `id`, `batteryLevel` (0–100) and `carrying` (boolean). Write a **constructor** that sets `id` from a parameter, sets `batteryLevel` to 100 and sets `carrying` to `false`. [5] (AO3)

(b) Add two methods to `Robot`: - `pickUp()` — if the robot is **not** already carrying, set `carrying` to `true` and return `true`; otherwise return `false`. - `drain(amount)` — reduce `batteryLevel` by `amount`, but never below 0. [5] (AO3)

(c) A class `ChargingRobot` **inherits** from `Robot` and adds a method `recharge()` that sets `batteryLevel` back to 100. Write `ChargingRobot` so that it inherits from `Robot` and adds this method. [2] (AO3)

(d) State the name of the OOP relationship between `ChargingRobot` and `Robot`, and give **one** advantage to the developers of using inheritance here. [2] (AO2)

Running total: 47

Question 5 — Computational methods [11]

(a) GridStore's route planner must plan around shelves, robots and walls. Explain what is meant by a **problem being solvable but intractable**, and give **one** example of a problem that is computable in principle but **not** in reasonable time for large inputs. [3] (AO2)

(b) Explain the difference between an **algorithm that guarantees the optimal solution** and a **heuristic**, and give **one** reason a heuristic could be preferred for GridStore's real-time route planning. [3] (AO2)

(c) The route planner must satisfy several conditions at once: each cell visited must be free, the path must not exceed the robot's remaining battery, and no two robots may occupy a cell at the same time-step. Name the computational method that describes solving a problem subject to a set of such restrictions. [1] (AO1)

(d) Explain what is meant by a **divide-and-conquer** approach, and state **one** standard algorithm (other than a sorting algorithm) that uses it. [2] (AO2)

(e) State **one** benefit of using **caching** of previously planned routes between common start and end cells, and **one** situation in which the cached route could become invalid. [2] (AO2)

Running total: 58

Question 6 — Big O analysis [11]

(a) Explain what **Big O notation** describes, and state why it expresses the **worst case** as a function of input size n rather than an exact running time in seconds. [2] (AO2)

(b) State the **worst-case** time complexity, using Big O, of each of the following: - (i) linear search of an unsorted list of n items - (ii) binary search of a sorted list of n items - (iii) merge sort of n items - (iv) accessing element i of an array by its index [4] (AO1)

(c) Consider the algorithm below.

```
i = n
while i > 1
    print(i)
    i = i DIV 2
endwhile
```

State its time complexity in Big O and **justify** your answer. [3] (AO2)

(d) Algorithm X runs in $O(n^2)$ and algorithm Y runs in $O(2^n)$. State which scales worse for large n , and explain why. [2] (AO2)

Running total: 69

Question 7 — Searching algorithms [10]

(a) State **one** precondition that must hold before a **binary search** can be used. [1] (AO1)

(b) Trace a **binary search** for the value **16** in the sorted list below (indices 0–8). At each step show `low`, `high`, `mid` and `list[mid]`, and state the total number of comparisons of `list[mid]` made. Use `mid = (low + high) DIV 2`. [5] (AO2)

Index:	0	1	2	3	4	5	6	7	8
Value:	4	9	16	23	31	38	52	61	77

(c) Complete the **binary search** function below so that it returns the **index** of `target` if found, or `-1` if not found. [4] (AO3)

```
function binarySearch(list, target)
    low = 0
    high = list.length - 1
    while low <= high
        mid = (low + high) DIV 2
        // ... complete: three branches ...
    endwhile
    // ... complete ...
endfunction
```

Running total: 79

Question 8 — Sorting algorithms [13]

- (a) A list contains [6, 3, 8, 2, 5]. Perform a **bubble sort**, writing out the **full state of the list after each complete pass**. State the pass on which your implementation could first detect that the list is sorted (no swaps). [5] (AO2)
- (b) State the **worst-case** and **best-case** time complexity of bubble sort (assuming a swap-detection optimisation), and describe the input that produces the **best** case. [3] (AO2)
- (c) **Quick sort** sorts the list [7, 2, 9, 4, 3, 8, 1] using the **last element of each (sub)list as the pivot**. Show the partitioning at each level: for each partition state the pivot, the resulting arrangement, and the left/right sublists produced, until the list is fully sorted. [3] (AO3)
- (d) Quick sort has an average-case complexity of $O(n \log n)$ but a worst case of $O(n^2)$. Describe the kind of input/pivot choice that causes quick sort's **worst case**, and state **one** advantage quick sort has over merge sort. [2] (AO2)

Running total: 92

Question 9 — Data structure algorithms [12]

- (a) A **circular queue** is implemented in an array `q` of size **5** (indices 0–4) with pointers `front` and `rear` and a `size` counter. Complete the `enqueue` procedure and the `dequeue` function below, including handling of the **full** and **empty** conditions. [5] (AO3)

```
procedure enqueue(value)
    if size == 5 then
        print("Queue full")
    else
        // ... complete: advance rear with wrap-around, store value, update size ...
    endif
endprocedure

function dequeue()
    if size == 0 then
        return "Queue empty"
    else
        // ... complete: read from front, advance front with wrap-around, update size, return
    endif
endfunction
```

- (b) Starting from an **empty** circular queue of size 5 (`front = 0`, `rear = -1`, `size = 0`), the following operations are performed in order:

```
enqueue(A), enqueue(B), enqueue(C), dequeue(), enqueue(D), dequeue(), enqueue(E)
```

State, **after all operations**, the value of `front`, the value of `rear`, and the values currently stored in the queue in order from front to rear. [3] (AO2)

(c) A **singly linked list** stores the nodes below. `head` points to the node at index 2. Each node holds (data, pointer), where the pointer is the index of the next node and `pointer = null` marks the end.

Index:	1	2	3	4
Data:	"Oslo"	"Lima"	"Cairo"	"Quito"
Pointer:	4	3	1	null

(i) State the order in which the data is read when the list is **traversed** from `head`. [2] (AO2) (ii) The node "Cairo" (index 3) is to be **deleted** from the list. State which single pointer changes and its new value, and give the new traversal order. [2] (AO3)

Running total: 107

Question 10 — Trees and traversals [11]

The values 55, 40, 65, 25, 48, 60, 80, 45 are inserted, in that order, into an initially empty **binary search tree (BST)**.

(a) Draw the resulting BST. [3] (AO3)

(b) Give the **pre-order**, **in-order** and **post-order** traversal output of your tree. [3] (AO2)

(c) A program needs to output the stored values in **descending** order. Describe how a traversal of the BST could be modified to achieve this directly, without sorting the output. [2] (AO2)

(d) State the **best-case** and **worst-case** time complexity of searching for a value in a BST, and describe the **shape** of tree that gives each case. [3] (AO2)

Running total: 118

Question 11 — Shortest path (Dijkstra / A*) [12]

Consider the weighted, **undirected** graph below.

Edges (undirected, weighted):
A-B = 4, A-C = 3, B-C = 1, B-D = 2, C-D = 5,
C-E = 8, D-E = 2, D-F = 6, E-F = 3

(a) Using **Dijkstra's algorithm**, find the shortest distance **and** the shortest **path** from **A** to **F**. Show a table of tentative distances and state clearly the order in which nodes are finalised (visited). [7] (AO3)

(b) **A*** orders nodes using cost-so-far plus a heuristic estimate of the remaining distance. State the formula **A*** uses, defining each term, and explain **one** advantage **A*** has over Dijkstra's algorithm when a good heuristic is available. [3] (AO2)

(c) State the property a heuristic must have for **A*** to be guaranteed to find the **optimal** path, and explain briefly what would go wrong if this property did not hold. [2] (AO2)

Running total: 130

Question 12 — Design and evaluation (extended response) [13]

GridStore (from Question 1) must store the live status of **every robot** (location, battery level, whether it is carrying an item, whether it is free) and, for each new order, repeatedly find the **nearest free robot** with **sufficient battery** to reach the pickup shelf. Robots change status many times per second, and thousands of orders arrive per minute.

Discuss how you would **design** the data storage and **searching** strategy for this problem. Your answer should:

- identify **suitable data structure(s)** for storing robot status and justify the choice, considering that status updates are very frequent;
- describe a **searching / algorithmic approach** for finding the nearest suitable robot, referring to relevant computational methods and/or graph/search techniques (e.g. spatial indexing, Dijkstra/A*, heuristics);
- analyse the **time-complexity** implications of your approach as the number of robots grows;
- **evaluate** at least one **trade-off** in your design (e.g. speed vs memory, accuracy vs responsiveness, simplicity vs scalability) and reach a justified conclusion.

The quality of your written communication will be assessed. **[13] (AO3)**

*Running total: 140***

Mark scheme

Marking guidance: award marks for valid alternative wording that conveys the same point. Code answers may be in OCR Exam Reference Language or any high-level language; award marks for correct logic regardless of syntax dialect. Trace marks require working to be shown, not just a final answer.

Question 1 — Computational thinking [11]

(a) [3]: - Technique = **abstraction** (specifically representational abstraction / modelling). [1] - The warehouse is represented as a grid of free/blocked cells, removing irrelevant detail (shelf height, robot colour). [1] - Why helpful: it produces a simpler model that is sufficient to plan collision-free routes and reduces the amount of data/processing, making path-finding feasible in real time. [1]

(b) [4] — two valid sub-problems, each with a sensible input and output (2 marks each): - e.g. *Allocate an order to a robot*: input = order details / item locations + robot statuses; output = chosen robot ID. [2] - e.g. *Plan a route for the robot*: input = robot start cell + target shelf cell + grid; output = sequence of cells / moves. [2] - (Other valid: locate the item on the shelf; check item availability; route robot to packing station; detect/avoid collisions.)

(c) [2]: - Concurrency must be managed because multiple robots act on / compete for the same shared resource (a cell) at the same time. [1] - Consequence if unmanaged: a **collision** / both robots try to occupy the same cell — a race condition leading to deadlock, damage or an inconsistent floor state. [1]

(d) [2] — any one well-explained reason for 2, or 1+1: - Categorising jobs lets the controller apply the **same handling rule to a whole class** of jobs (e.g. always service urgent before standard) [1], giving predictable, consistent prioritisation without examining every job individually, which helps guarantee response times. [1]

Question 2 — Programming techniques [12]

(a) [1] — e.g. an iterative solution uses a loop and fixed memory; a recursive solution calls itself and uses the call stack (extra memory per call). [1] (any one clear difference)

(b) [5] — Output, in order:

```
[15, 25, 35]
```

```
3
```

- `nums` is passed **by reference**, so `adjust` modifies the original array; each element gains 5 → `[15, 25, 35]`. [1] (effect on array) + [1] (correct printed array)
- `step` is passed **by value**, so the value 5 is used inside the subroutine; nothing about a copy affects the result here but the value added is 5. [1]
- `count` is a **global** variable, in scope inside the procedure; it is incremented once per element (3 elements) → prints `3`. [1]
- Correct ordering of both printed lines. [1]

(c) [3]: - Scope = the region of a program in which a variable is **accessible / has meaning / exists**. [1] - (Local = within the subroutine declared; global = whole program.) [1] - Disadvantage of a global: any part of the program can change it, causing **unintended side-effects / hard-to-trace bugs**, and it reduces reusability of subroutines / risks name clashes. [1] (any one)

(d) [3]:

```
total = 0
k = 1
while k <= 5
    total = total + k
    k = k + 1
endwhile
```

- Correct initialisation of `total` and `k` before the loop. [1]
- Correct loop condition `while k <= 5`. [1]
- Body adds `k` to `total` and increments `k` inside the loop (so it terminates). [1]

Question 3 — Recursion and the call stack [10]

(a) [2]: - Base case = `if exp == 0 then return 1` (stops the recursion). [1] - Recursive/general case = `return base * power(base, exp - 1)` (calls itself with `exp` reduced, moving towards the base case). [1]

(b) [5] — `power(3, 4)` . Frames pushed (each waits for the one below):

```
power(3,4) -> 3 * power(3,3)
power(3,3) -> 3 * power(3,2)
power(3,2) -> 3 * power(3,1)
power(3,1) -> 3 * power(3,0)
power(3,0) -> base case, returns 1
```

Unwinding (popping):

Frame	Evaluates	Returns
power(3,4)	3 * 27	81
power(3,3)	3 * 9	27
power(3,2)	3 * 3	9
power(3,1)	3 * 1	3
power(3,0)	base case	1

Marks: - Correct chain of calls pushed: `exp = 4 → 3 → 2 → 1 → 0`. [2] - Base case `power(3,0) = 1` identified. [1] - Correct returns on unwinding: `1 → 3 → 9 → 27`. [1] - Final frame returns `3 × 27 = 81`. [1]

(c) [1] — `power(3, 4) = 81`. [1]

(d) [2]: - Without the base case, `exp` keeps decreasing (4, 3, 2, 1, 0, -1, -2 ...) and the function calls itself forever / never returns. [1] - The call stack keeps growing until memory is exhausted: a **stack overflow** error. [1]

Question 4 — Object-oriented programming [14]

(a) [5] — e.g.:

```
class Robot
  private id
  private batteryLevel
  private carrying

  public procedure new(robotID)
    id = robotID
    batteryLevel = 100
    carrying = false
  endprocedure
endclass
```

- Class declared with name `Robot`. [1]
- Three attributes declared `private`. [1]
- Constructor takes an `id` parameter and assigns it to the attribute. [1]
- Constructor sets `batteryLevel = 100`. [1]
- Constructor sets `carrying = false`. [1]

(b) [5]:

```
public function pickUp()
  if carrying == false then
    carrying = true
    return true
  else
    return false
  endif
endfunction

public procedure drain(amount)
  batteryLevel = batteryLevel - amount
  if batteryLevel < 0 then
    batteryLevel = 0
  endif
endprocedure
```

- `pickUp` checks `carrying == false`. [1]
- `pickUp` sets `carrying = true` and returns `true` when not carrying. [1]
- `pickUp` returns `false` when already carrying. [1]
- `drain` subtracts `amount` from `batteryLevel`. [1]

- `drain clamps batteryLevel` so it is never below 0. [1]

(c) [2]:

```
class ChargingRobot inherits Robot
  public procedure recharge()
    batteryLevel = 100
  endprocedure
endclass
```

- Correct inheritance syntax (`inherits Robot` / `extends`). [1]
- Adds a `recharge()` method that sets `batteryLevel = 100`. [1] (Accept use of a setter / `super` if `batteryLevel` is treated as needing access via the parent.)

(d) [2]: - Relationship = **inheritance** ("is-a": a ChargingRobot is a Robot). [1] - Advantage: code reuse — ChargingRobot inherits all the attributes/methods of Robot without rewriting them / changes to shared behaviour need only be made once. [1] (any one)

Question 5 — Computational methods [11]

(a) [3]: - A problem is **solvable but intractable** if an algorithm exists to solve it exactly, but the time/resources required grow so fast (e.g. exponentially/factorially) that it cannot be solved in a reasonable time for large inputs. [1] (idea of solvable) + [1] (idea of not in reasonable time / explosive growth) - Example: the **Travelling Salesman Problem** / optimal scheduling / optimal multi-robot routing visiting all points. [1]

(b) [3]: - Optimal algorithm guarantees the best possible solution but may be slow/exhaustive. [1] - Heuristic uses a "rule of thumb" to find a **good-enough** solution quickly without guaranteeing optimality. [1] - Reason preferred for GridStore: routes must be planned in **real time** for many robots; a near-optimal route found instantly is more useful than a perfect route found too late. [1]

(c) [1]: - **Constraint-based / constraint satisfaction** problem solving (allow "backtracking" as the technique used to solve it). [1]

(d) [2]: - Divide and conquer = break a problem into smaller sub-problems, solve each (often recursively), then **combine** the results. [1] - Example (non-sorting): **binary search** (or the recursive nearest-pair, fast Fourier transform, etc.). [1]

(e) [2]: - Benefit of caching: a previously planned route between common cells can be **reused instantly**, saving repeated computation / reducing latency. [1] - Becomes invalid when: the grid changes — e.g. a shelf/robot now blocks a cell on the cached route, or a cell becomes congested — so the stored route is no longer free/valid. [1]

Question 6 — Big O analysis [11]

(a) [2]: - Big O describes how an algorithm's **run-time (or space) requirement grows as the input size n grows** / its order of growth in the worst case. [1] - It uses n rather than seconds because actual time depends on hardware/implementation; Big O captures the **rate of growth**, which is what matters for scalability and lets algorithms be compared independently of the machine. [1]

(b) [4] — 1 mark each: - (i) Linear search (unsorted): **$O(n)$** . [1] - (ii) Binary search: **$O(\log n)$** . [1] - (iii) Merge sort: **$O(n \log n)$** . [1] - (iv) Array access by index: **$O(1)$** . [1]

(c) [3]: - Complexity = **$O(\log n)$** . [1] - Justification: `i` is **halved** each iteration (`i = i DIV 2`), so the number of iterations is roughly $\log_2 n$ before `i` reaches 1. [1] - The loop body is constant work, so the total grows logarithmically with `n`. [1]

(d) [2]: - **$Y (O(2^n))$** scales worse. [1] - Exponential growth (2^n) eventually overtakes any polynomial (n^2) and grows far faster as `n` increases — e.g. each extra input item doubles the work for `Y`, whereas `X` grows only quadratically. [1]

Question 7 — Searching algorithms [10]

(a) [1] — The list must be **sorted** (in order). [1] (Accept: random/direct access to elements by index is available.)

(b) [5] — Searching for 16 in `[4, 9, 16, 23, 31, 38, 52, 61, 77]` (indices 0–8), `mid = (low+high) DIV 2`:

Step	low	high	mid	list[mid]	Action
1	0	8	4	31	<code>16 < 31</code> → search left, <code>high = 3</code>
2	0	3	1	9	<code>16 > 9</code> → search right, <code>low = 2</code>
3	2	3	2	16	found

Marks: - Step 1 correct (`mid = 4`, `list[mid] = 31`, go left, `high = 3`). [1] - Step 2 correct (`mid = 1`, `list[mid] = 9`, go right, `low = 2`). [1] - Step 3 correct (`mid = 2`, `list[mid] = 16`, found). [1] - `mid` calculated correctly each step using integer division. [1] - Number of comparisons of `list[mid]` = **3**. [1]

(c) [4]:

```
function binarySearch(list, target)
  low = 0
  high = list.length - 1
  while low <= high
    mid = (low + high) DIV 2
    if list[mid] == target then
      return mid
    elseif list[mid] < target then
      low = mid + 1
    else
      high = mid - 1
    endif
  endwhile
  return -1
endfunction
```

- `if list[mid] == target then return mid`. [1]
- `list[mid] < target` branch sets `low = mid + 1`. [1]
- `else` branch sets `high = mid - 1`. [1]

- return -1 after the loop (not found). [1]

Question 8 — Sorting algorithms [13]

(a) [5] — Bubble sort of [6, 3, 8, 2, 5], full list after each complete pass:

Pass	List state after pass	Swaps?
1	[3, 6, 2, 5, 8]	yes
2	[3, 2, 5, 6, 8]	yes
3	[2, 3, 5, 6, 8]	yes
4	[2, 3, 5, 6, 8]	no swaps

- [3, 6, 2, 5, 8] after pass 1. [1]
- [3, 2, 5, 6, 8] after pass 2. [1]
- [2, 3, 5, 6, 8] after pass 3 (now sorted). [1]
- Pass 4 makes no swaps. [1]
- States that sorting can first be detected (no swaps) on **pass 4**. [1] (Accept detection on pass 3 if the candidate's implementation tracks "no swap" within the same pass that produces the sorted order; award the mark for a consistent, correctly-justified pass number.)

(b) [3]: - Worst case: $O(n^2)$. [1] - Best case: $O(n)$ (with swap-detection). [1] - Best case occurs when the input is **already sorted** — one pass with no swaps detects completion. [1]

(c) [3] — Quick sort of [7, 2, 9, 4, 3, 8, 1], pivot = last element of each (sub)list.

- Partition whole list, pivot = 1: everything ≥ 1 goes right $\rightarrow$ [1 | 2, 9, 4, 3, 8, 7]. Left = [], right = [2, 9, 4, 3, 8, 7]. [1]
- Partition [2, 9, 4, 3, 8, 7], pivot = 7 $\rightarrow$ [2, 4, 3 | 7 | 9, 8]. Left = [2, 4, 3], right = [9, 8]. [1]
- Partition [2, 4, 3], pivot = 3 $\rightarrow$ [2 | 3 | 4]; partition [9, 8], pivot = 8 $\rightarrow$ [8 | 9].
Recombining gives the sorted list [1, 2, 3, 4, 7, 8, 9]. [1]

(Full picture for reference:)

```

[7, 2, 9, 4, 3, 8, 1]      pivot 1 -> 1 | [2,9,4,3,8,7]
      [2, 9, 4, 3, 8, 7]  pivot 7 -> [2,4,3] | 7 | [9,8]
    [2,4,3]              pivot 3 -> [2] | 3 | [4]
          [9,8]          pivot 8 -> [8] | 9
Result: [1, 2, 3, 4, 7, 8, 9]

```

(d) [2]: - Worst case occurs when the pivot is consistently the smallest or largest element (e.g. an **already-sorted (or reverse-sorted) list with last-element pivot**), giving maximally unbalanced partitions of size $n-1$ and $0 \rightarrow O(n^2)$. [1] - Advantage over merge sort: quick sort sorts **in place**, needing only $O(\log n)$ extra stack space rather than $O(n)$ auxiliary memory. [1] (Accept: typically faster in practice due to good cache behaviour / lower constant factors.)

Question 9 — Data structure algorithms [12]

(a) [5]:

```
procedure enqueue(value)
  if size == 5 then
    print("Queue full")
  else
    rear = (rear + 1) MOD 5
    q[rear] = value
    size = size + 1
  endif
endprocedure

function dequeue()
  if size == 0 then
    return "Queue empty"
  else
    value = q[front]
    front = (front + 1) MOD 5
    size = size - 1
    return value
  endif
endfunction
```

- enqueue: advance rear with wrap-around $(\text{rear} + 1) \text{ MOD } 5$. [1]
- enqueue: store $q[\text{rear}] = \text{value}$ and increment size. [1]
- dequeue: read $\text{value} = q[\text{front}]$ before advancing. [1]
- dequeue: advance front with wrap-around $(\text{front} + 1) \text{ MOD } 5$ and decrement size. [1]
- dequeue: returns the dequeued value. [1]

(b) [3] — trace from empty ($\text{front} = 0$, $\text{rear} = -1$, $\text{size} = 0$):

Operation	front	rear	size	contents (front → rear)
enqueue(A)	0	0	1	A
enqueue(B)	0	1	2	A, B
enqueue(C)	0	2	3	A, B, C
dequeue()	1	2	2	B, C
enqueue(D)	1	3	3	B, C, D
dequeue()	2	3	2	C, D
enqueue(E)	2	4	3	C, D, E


- Final front = 2. [1]
- Final rear = 4. [1]
- Contents front → rear = C, D, E. [1]

(c)(i) [2] — Traverse from head (index 2): - 2 ("Lima") → 3 ("Cairo") → 1 ("Oslo") → 4 ("Quito"). **[1]** - Output order: **Lima, Cairo, Oslo, Quito. [1]**

(c)(ii) [2] — Delete "Cairo" (index 3). Node 2 ("Lima") currently points to 3; "Cairo" points to 1: - Node 2 ("Lima") pointer changes from 3 to **1** (bypasses Cairo). **[1]** - New traversal: **Lima, Oslo, Quito. [1]**

Question 10 — Trees and traversals [11]

(a) [3] — Insert order 55, 40, 65, 25, 48, 60, 80, 45:



- Root 55 with 40 (left) and 65 (right). **[1]**
- 25 and 48 as children of 40; 60 and 80 as children of 65. **[1]**
- 45 placed as **left child of 48** ($45 > 40 \rightarrow$ right to 48; $45 < 48 \rightarrow$ left of 48). **[1]**

(b) [3]: - **Pre-order** (root, left, right): 55, 40, 25, 48, 45, 65, 60, 80. **[1]** - **In-order** (left, root, right): 25, 40, 45, 48, 55, 60, 65, 80. **[1]** - **Post-order** (left, right, root): 25, 45, 48, 40, 60, 80, 65, 55. **[1]**

(c) [2]: - A standard in-order traversal visits left, root, right and gives **ascending** order. To get **descending** order, visit **right subtree, then root, then left subtree** (a "reverse in-order" traversal). **[1]** - This outputs the values largest-first directly, with no separate sort needed. **[1]**

(d) [3]: - Best case = **$O(\log n)$** . **[1]** — when the tree is **balanced**, each comparison halves the remaining nodes. **[1]** - Worst case = **$O(n)$** , when the tree is **unbalanced / degenerate** (effectively a linked list, e.g. from inserting already-sorted data). **[1]**

Question 11 — Shortest path (Dijkstra / A*) [12]

Graph edges (undirected): A-B=4, A-C=3, B-C=1, B-D=2, C-D=5, C-E=8, D-E=2, D-F=6, E-F=3.

(a) [7] — Dijkstra from A. Tentative distances (∞ = unknown); each round finalise the unvisited node with the smallest tentative distance and relax its neighbours.

Finalised this round	A	B	C	D	E	F
start	0	∞	∞	∞	∞	∞
A (0)	0	4 (via A)	3 (via A)	∞	∞	∞
C (3)	0	4 (A) / 4 (via C: 3+1)	3	8 (via C: 3+5)	11 (via C: 3+8)	∞
B (4)	0	4	3	6 (via B: 4+2)	11	∞
D (6)	0	4	3	6	8 (via D: 6+2)	12 (via D: 6+6)
E (8)	0	4	3	6	8	11 (via E: 8+3)
F (11)	0	4	3	6	8	11

Order finalised: **A, C, B, D, E, F.**

- Shortest distance A $\rightarrow$ F = **11**. [1]
- Shortest path = **A $\rightarrow$ B $\rightarrow$ D $\rightarrow$ E $\rightarrow$ F**. [1]

(Path note: F's best is via E (8+3=11). E's best is via D (6+2=8). D's best is via B (4+2=6). B's best is via A (4). C, although finalised earlier at 3, does not lie on the shortest path to F because routing through B gives a cheaper D.)

Working marks: - Initialise A = 0, others ∞ ; relax A's neighbours (B = 4, C = 3). [1] - Finalise C (= 3); relax to give D = 8 and E = 11 (B stays 4 since 3+1 = 4 is not smaller). [1] - Finalise B (= 4); update D to 6 (4+2). [1] - Finalise D (= 6); update E to 8 (6+2) and F to 12 (6+6). [1] - Finalise E (= 8); update F to 11 (8+3); correct visiting order A, C, B, D, E, F shown. [1]

(Award the distance and path marks (the first two listed) plus the five working marks = 7.)

(b) [3]: - A* orders nodes by **$f(n) = g(n) + h(n)$** . [1] - where **$g(n)$** = the actual cost of the path found so far from the start to n, and **$h(n)$** = the heuristic estimate of the remaining cost from n to the goal. [1] - Advantage: the heuristic **guides** the search towards the goal, so A* typically expands **fewer nodes** than Dijkstra (which expands uniformly outward), finding the path faster. [1]

(c) [2]: - The heuristic must be **admissible** — it must **never overestimate** the true remaining cost to the goal. [1] - If it overestimated, A* could discard the genuinely shortest route (treating it as more expensive than it is) and return a **sub-optimal path**. [1] (Accept "consistent/monotonic" as the property.)

Question 12 — Design and evaluation [13]

Level-of-response marking (0–13). Award the band that best fits, then a mark within it.

- **Level 3 (10–13):** A well-structured response that addresses **all four** required points. Data structure choice is justified (with the frequent-update requirement considered), the searching approach is appropriate and explained with reference to relevant computational methods/techniques, time complexity is analysed correctly as the number of robots grows, and at least one trade-off is evaluated with a reasoned conclusion. Technically accurate, correct terminology throughout. QWC strong.
- **Level 2 (5–9):** Addresses most points with some justification. Data structure and search approach identified and mostly appropriate; some complexity discussion; a trade-off mentioned but evaluation limited. Mostly accurate, reasonable structure.

- **Level 1 (1–4):** Fragmentary/descriptive. Names a structure and/or search method with little justification; little or no complexity analysis or evaluation. Some inaccuracies.
- **0 marks:** nothing creditworthy.

Indicative content (credit any valid, well-argued combination):

Data structures: - A **hash table / dictionary keyed by robot ID** giving **$O(1)$** average-case status updates and lookups — important because robots change status many times per second. - A **spatial structure** (uniform grid of buckets keyed by floor cell, a quadtree, or a k-d tree) so "nearest free robot" queries only examine robots near the pickup cell rather than all robots. - A **priority queue / min-heap** to retrieve robots ordered by distance/ETA, or to drive Dijkstra/A *expansion*. - *Optionally model the warehouse floor as a weighted graph** of cells for true travel-distance nearest-robot queries.

Searching / algorithmic approach: - For straight-line nearest: query the spatial index for candidate robots in the pickup cell and adjacent cells, filter to those that are **free and have sufficient battery**, then take the minimum distance (a localised linear scan of few candidates). - For travel-distance nearest on the grid: **Dijkstra / A*** from the pickup cell; A* with an admissible straight-line (Manhattan/Euclidean) heuristic explores fewer cells. - A **heuristic** (e.g. straight-line distance, ignoring obstacles) may be "good enough" for real-time selection rather than computing exact path length to every robot.

Time complexity: - Naïve scan of all robots = **$O(n)$** per order → with thousands of orders/min and many robots this becomes a bottleneck. - A spatial index reduces each query to robots in nearby cells, **$\sim O(k)$** with $k \ll n$, or **$O(\log n)$** for tree structures — far more scalable as n grows. - Status updates should be **$O(1)$** (hash table) or $O(\log n)$ so very frequent changes don't dominate.

Trade-offs to evaluate (need at least one, with a judgement): - **Speed vs accuracy:** a straight-line heuristic is fast but ignores blocked cells/congestion; exact path-finding is accurate but costly. - **Speed vs memory:** spatial indexes/caches consume extra memory to gain query speed. - **Responsiveness vs optimality:** in real time, dispatching a near-optimal robot instantly beats finding the perfect robot too late. - **Concurrency:** updates and queries happen simultaneously — must avoid race conditions when marking a robot as taken (links to Q1c).

A response reaching Level 3 makes a clear recommendation (e.g. hash table keyed by robot ID for $O(1)$ status + spatial grid for candidate selection + A* for travel-distance ranking) and justifies it against these trade-offs.

Verified mark total

Question	Marks
Q1 Computational thinking	11
Q2 Programming techniques	12
Q3 Recursion & call stack	10
Q4 Object-oriented programming	14
Q5 Computational methods	11
Q6 Big O analysis	11
Q7 Searching algorithms	10
Q8 Sorting algorithms	13
Q9 Data structure algorithms	12
Q10 Trees and traversals	11
Q11 Shortest path (Dijkstra / A*)	12
Q12 Design and evaluation	13
TOTAL	140

Verified total = 140 marks.