

Knowledge Organiser — 2.3 Algorithms

OCR A Level Computer Science H446, Component 02, Section 2.3.1 — Algorithms.

Big O complexity classes

Class	Name	Growth as n increases	Example
O(1)	Constant	No change with n	<code>array[i]</code> ; stack push/pop
O(log n)	Logarithmic	Doubling n adds one step	Binary search
O(n)	Linear	Doubling n doubles time	Linear search; single loop
O(n log n)	Linearithmic	Slightly worse than linear	Merge/quick sort (average)
O(n²)	Polynomial (quadratic)	Doubling n quadruples time	Bubble/insertion sort; nested loops
O(2ⁿ)	Exponential	One extra item doubles time	Naïve recursive Fibonacci; brute-force TSP

- Big O = **worst-case** growth of **time** (operations) or **space** (memory) vs input size n.
- Drop constants & lower-order terms: $3n^2 + 5n + 17 \rightarrow \mathbf{O(n^2)}$ (dominant term wins).
- **Tractable** = solvable in polynomial time or better; **intractable** = only exponential ($O(2^n)$) or worse.

Standard algorithm complexities (VERIFIED)

Algorithm	Best	Average	Worst	Space	Stable?
Linear search	O(1)	O(n)	O(n)	O(1)	—
Binary search	O(1)	O(log n)	O(log n)	O(1) iter / O(log n) rec	—
Bubble sort	O(n)*	O(n ²)	O(n ²)	O(1)	Yes
Insertion sort	O(n)	O(n ²)	O(n ²)	O(1)	Yes
Merge sort	O(n log n)	O(n log n)	O(n log n)	O(n)	Yes
Quick sort	O(n log n)	O(n log n)	O(n²)	O(log n)	No

* Bubble best case O(n) only **with a swap flag** on already-sorted data; otherwise O(n²).

How each one works (one line)

Algorithm	How it works
Linear search	Check each element in turn from the start until found or list ends. Works on unsorted data.
Binary search	Check middle of a sorted list; discard the half that cannot contain target; repeat (halves space each pass).
Bubble sort	Repeatedly compare adjacent pairs and swap if out of order; largest "bubbles" to end each pass.
Insertion sort	Build a sorted front section; insert each next element into place by shifting larger ones right.
Merge sort	Divide-and-conquer: recursively split in half until single items, then merge sublists in order.
Quick sort	Divide-and-conquer: pick a pivot, partition (smaller left / larger right), recursively sort each side.

Binary search precondition: the list MUST be sorted. Quote this in every answer.

Data structure operation complexities

Structure	Operations	Complexity
Stack (LIFO)	push / pop / peek / isEmpty	$O(1)$
Queue (FIFO)	enqueue (rear) / dequeue (front)	$O(1)$
Array	access by index	$O(1)$; insert/delete shifts $\rightarrow O(n)$
Linked list	traverse / search	$O(n)$; insert/delete given node $\rightarrow O(1)$ (re-point)
Binary search tree	search / insert	$O(\log n)$ avg (balanced), $O(n)$ worst (degenerate); traversal $O(n)$

- Circular queue uses **MOD** to wrap pointers and reuse freed space.
- **DFS uses a stack** (or recursion); **BFS uses a queue**.

Tree traversals

Example tree: root 50; 30 (L) with children 20,40; 70 (R) with children 60,80.

Traversal	Order	Output	Use
Pre-order	Node $\rightarrow$ Left $\rightarrow$ Right	50, 30, 20, 40, 70, 60, 80	Copy a tree; prefix expressions
In-order	Left $\rightarrow$ Node $\rightarrow$ Right	20, 30, 40, 50, 60, 70, 80	Ascending sorted output from a BST
Post-order	Left $\rightarrow$ Right $\rightarrow$ Node	20, 40, 30, 60, 80, 70, 50	Delete a tree; evaluate RPN (postfix)
Breadth-first (BFS)	Level by level, L $\rightarrow$ R (uses a queue)	50, 30, 70, 20, 40, 60, 80	Shortest path in unweighted graph

- Mnemonic — position of **Node**: **Pre** = first, **In** = middle, **Post** = last.
- Post-order is the **depth-first** traversal the spec names explicitly.

Path-finding: Dijkstra vs A*

	Dijkstra	A*
Cost function	$g(n)$ only	$f(n) = g(n) + h(n)$
Heuristic?	No (uninformed)	Yes (informed)
Search shape	Outward in all directions	Directed toward the goal
Typical use	Shortest path to all nodes	Shortest path to one goal
Efficiency	Explores more nodes	Usually explores fewer → faster
Optimality	Optimal (non-negative weights)	Optimal if heuristic admissible

- $g(n)$ = actual cost from start to n . $h(n)$ = heuristic *estimate* of cost from n to goal (e.g. straight-line/Manhattan). $f(n)$ = estimated total cost; always expand lowest $f(n)$ next.
- **A* adds a heuristic estimate of the remaining distance to the goal** — that is the key difference.
- Heuristic must be **admissible** (never overestimates) for A* to guarantee the optimal path. If $h(n)=0$, A* reduces exactly to Dijkstra.
- Dijkstra: always visit the unvisited node with smallest tentative distance; relax neighbours; once visited, distance is final (don't update settled nodes).

Must-know definitions

Term	Definition
Algorithm	A precise, finite sequence of steps to solve a problem.
Big O notation	How an algorithm's time/space requirement grows with input size n (worst case).
Time complexity	Number of operations as a function of input size.
Space complexity	Memory required as a function of input size.
Tractable / Intractable	Solvable in polynomial time or better / only exponential or worse.
Divide and conquer	Break into subproblems, solve, combine (merge/quick sort, binary search).
Pivot	Element chosen in quick sort to partition the list.
Stable sort	Preserves relative order of equal elements.
Heuristic	A rule-of-thumb estimate; in A*, estimate $h(n)$ of remaining cost to goal.
Admissible heuristic	One that never overestimates the true remaining cost → guarantees A* optimality.
Traversal	Visiting every node of a structure in a defined order.

Top exam reminders

- **Binary search is $O(\log n)$ and REQUIRES sorted data** — always state the precondition.

- **Quick sort worst case is $O(n^2)$** (poor pivot / already-sorted); **merge sort guarantees $O(n \log n)$** but uses **$O(n)$ extra space** — classic time-vs-space trade-off.
- **Always justify** a Big O, e.g. " $O(n^2)$ because of the nested loop — each of n passes makes up to n comparisons."
- **DFS uses a stack; BFS uses a queue.** Don't mix them up.
- **In-order traversal of a BST outputs values in ascending order**; bubble sort best case is $O(n)$ only with a swap flag.